

# Introduction to Machine Learning (by Implementation)

## Lecture 2: Gradient Descent

Ian J. Watson

University of Seoul

University of Seoul Graduate Course 2019



**KRF** KOREA RESEARCH FELLOWSHIP  
해외 우수신진연구자 유치사업



# List comprehensions

- Review: it will often be useful to quickly make transformations from one list to another, applying some transformation to each element
- This can be done with a for loop and append:

```
oldList = [1, 2, 3, 4]
newList = []
for element in oldList:
    newList.append(element**2)
```

- But it can be done quicker with list comprehensions

```
newList = [element**2 for element in oldList]
```
- Use enumerate if you need an index

```
[element if index == 0 else 0 for (index, element) in enumerate(oldList)]
```
- If you don't need to use the element, can write `_`

```
[index for index, _ in enumerate(oldList)]
```

- So far, we've talked about python and done some introductory coding
- Today, we'll introduce and implement another concept that will be important going forward: gradient descent
- Often, we will have large complex models whose parameters we want to tune, in order to find the "best" model
  - "Best" will mean something like "minimizes the error of the model" or "maximizes the likelihood of the data"
- Therefore, we will need to a routine to find minima of functions
  - Generally, are models can be arbitrarily complex, and not amenable to analytical minimization (solving  $\nabla f = 0$  directly)
- Gradient descent is a method to minimize a function numerically (as opposed to analytically)

# Analogy: Steepest descent




- A climber is trying to find his way down a mountain in deep fog, how should he proceed?
- One idea is to try to always go downhill the fastest way possible
- So, he figures out which direction has the steepest descent (ie which way is downhill), then takes a step in that direction
- After the step, he checks again, and takes another step
- He keeps proceeding in this manner until he can't go downhill anymore, he's reached the bottom

# Gradient Descent

- From calculus,  $\nabla f(\mathbf{x})$  gives the direction of largest increase of  $f$  at  $\mathbf{x}$  (if its  $\mathbf{0}$ , we are at a minimum and done)
- Equivalently,  $-\nabla f(\mathbf{x})$  gives direction of largest decrease, so  $f(\mathbf{x} - \gamma \nabla f(\mathbf{x})) < f(\mathbf{x})$  (at least, for some  $\gamma$  small enough)
- We will define a sequence  $\mathbf{x}_i$  to find the minimum:
  - Start with some random position  $\mathbf{x}_0$
  - Iterate:
    - Find  $\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma_n \nabla f(\mathbf{x}_n)$
    - Stop if  $|f(\mathbf{x}_{n+1}) - f(\mathbf{x}_n)| < \epsilon$ , i.e. we're not reducing further, so we're close to the minimum
  - Return the final  $\mathbf{x}_n$
- $\gamma_n$  can be different for each iteration, we'll find the best  $\gamma_n$  by checking several possible values
- $\epsilon$  is the *tolerance*, how close to a minima do we need to be before stopping (we could also check  $|\nabla f(\mathbf{x}_n)| < \epsilon$ )

# Example function



- Shows how the algorithm picks out different paths depending on starting point
- Lines are contours of equal value

# Partial Derivative Estimation


- To run gradient descent, we must find the partial derivatives
- Maybe we are given only the function,  $f$ , without the analytic form, or the derivative functions, what can we do then?
- We will have to come up with a numerical estimate for the derivative
- The simplest way to do this is to simply take a partial difference quotient:

- $\frac{\partial f}{\partial x_i}(\mathbf{x}) \approx \frac{f(\mathbf{x} + h \cdot \mathbf{e}_i) - f(\mathbf{x})}{h}$

- $\mathbf{e}_i$  is a unit vector in the  $x_i$  direction

- As  $h \rightarrow 0$ , our estimate  $\rightarrow \frac{\partial f}{\partial x_i}(\mathbf{x})$



- There are better (= converge faster) methods, but this suffices



- Could take difference in forward or backward direction, or an average
  - Starting point for better estimates
- In limit as  $\Delta x \rightarrow 0$ , all the methods converge to the true derivative
  - So, just take the equation shown when building your code

# Possible Issues

- Multiple minima: gradient descent will find the closest minimum, this is not necessarily the global minimum
  - Can be dealt with by testing multiple starting points and using techniques like simulated annealing



- Gradient descent also slowly converges near the minimum, can take a long time to go from near the minimum to the true minimum
  - Possible to detect and methods to speed up conversion
- A real numerical minimizer (e.g. MINUIT, used in particle physics libraries like ROOT) will take these considerations into account



- We will deal with functions of several variables
- The way we will represent this python is as follows:
- A function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ , will be a python function `f`, which takes a list of  $n$  floats, and returns a float
  - $f(x, y) = x^2 + y^2 \rightarrow f = \text{lambda } x: x[0]**2 + x[1]**2$
- We can then represent the basis vectors like  $\mathbf{e}_i = [0, \dots, 1, \dots, 0]$  with the 1 in the  $i$ 'th position
  - E.g. for  $\mathbb{R}^2$ , the basis vectors are  $\mathbf{e}_0 = [1, 0]$  and  $\mathbf{e}_1 = [0, 1]$

# Exercises

We will break up the algorithm into pieces to make it easier to test one by one. The functions should use the previous functions

- `step(v: List[float], direction: List[float], step_size: float) -> List[float]` increments the point  $v$  in the direction by `step_size`, i.e.  $\text{step}(\mathbf{v}, \mathbf{d}, s) = \mathbf{v} + s \cdot \mathbf{d}$
- `move_point_along_ei(x: List[float], i: int, h: float) -> List[float]` takes the  $x$ , and shifts it by  $h \cdot \mathbf{e}_i$ , that is  $\mathbf{x} + h \cdot \mathbf{e}_i$
- `partial_difference_quotient(f, v, i, h)` which estimates the  $i$ 'th partial derivative of the function  $f$  at  $v$  with a step size of  $h$ ,  
$$\frac{f(\mathbf{x} + h \cdot \mathbf{e}_i) - f(\mathbf{x})}{h}$$
- `estimate_gradient(f: Fn, v: List[float], h=0.00001)` which uses `partial_difference_quotient` to estimate  $\nabla f$  at  $v$
- Check that your `estimate_gradient` function gives sensible results, add some tests to test `_gradient_descent.py`
  - E.g.  $f(x, y) = x^2 + y^2 \implies \nabla f(x, y) = (2x, 2y)$ , does your function give a result close to the real gradient at  $(x, y) = (0, 0)$ , and  $(x, y) = (1, 1)$ ?

# Exercises (minimize)

- Write a function `minimize(f, df, x0, step_sizes, tol)` which implements the gradient descent algorithm
  - `f: Callable[[List[float]], float]` is the function we want to minimize
    - `Callable[a, b]` is the type of a function that takes in `a` (list of types) and returns a type `b`
  - `df: Callable[[List[float]], List[float]]` should be the gradient function, i.e. it takes a point  $\mathbf{x}$  and returns  $\nabla f(\mathbf{x})$ 
    - If you don't have a function for the exact gradient, you can use your estimate by passing `lambda x: estimate_gradient(f, x)` instead of `df`
  - `x0: List[float]` the initial position, point to start the gradient descent minimization from
  - `step_sizes: List[float]`, this is the list we choose  $\gamma_n$  from. Test each element of the list, and choose the one that gives the lowest value of `f`. Use default value `[10, 1, 0.1, 0.01, 0.001]`
  - `tol: float` the tolerance, when `f` from one step to the next changes by less than this amount, end the search. Use default value `1e-5`

- Summary of the algorithm:
  - Find  $\nabla f$  at  $x_n$  using `df` (first time through it will be )
  - Set  $\mathbf{x}_{n+1} = \mathbf{x}_n + h \cdot \nabla f$ , where  $h$  is chosen from `step_sizes` to give the lowest value of the function
  - When when the function changes by less than `tol`, stop and return the previous estimate, not the new estimate for the step, otherwise loop
    - If  $|f(\mathbf{x}_{n+1}) - f(\mathbf{x}_n)| < \text{tol}$ , return  $x_n$
- Check the `test_gradient_descent.py` function `test_minimize` for examples of how the code should work
- Check that your `minimize` function finds a value close to the true minimum for some known functions, and add to the test file
  - Do you find  $(0,0)$  if you minimize  $f(x,y) = x^2 + y^2$ ? Does it change if you use the real  $\nabla f$  or your `estimate_gradient` version?